



# Chapter 8

## Sequential-Access Files

C++ How to Program,  
Late Objects Version, 7/e



## OBJECTIVES

In this chapter you'll learn:

- The data hierarchy from bits, to files to databases.
- To create, read, write and update sequential files.
- Some of the key streams that are associated with file processing.



- 8.1** Introduction
- 8.2** Data Hierarchy
- 8.3** Files and Streams
- 8.4** Creating a Sequential File
- 8.5** Reading Data from a Sequential File
- 8.6** Updating Sequential Files
- 8.7** Wrap-Up



# 8.1 Introduction

- ▶ Storage of data in memory is temporary.
- ▶ **Files** are used for **data persistence**—permanent retention of data.
- ▶ Computers store files on **secondary storage devices**, such as hard disks, CDs, DVDs, flash drives and tapes.
- ▶ In this chapter, we explain how to build C++ programs that create, update and process sequential files.
- ▶ We examine techniques for input of data from, and output of data to, **string** streams rather than files in Chapter 18, Class **string** and String Stream Processing.



## 8.2 Data Hierarchy

- ▶ Ultimately, all data items that digital computers process are reduced to combinations of zeros and ones.
  - It's simple and economical to build electronic devices that can assume two stable states—one state represents 0 and the other represents 1.
- ▶ The smallest data item that computers support is called a **bit**
  - Short for “**binary digit**”—a digit that can assume one of two values
  - Each data item, or bit, can assume either the value 0 or the value 1.
- ▶ Computer circuitry performs various simple bit manipulations, such as examining the value of a bit, setting the value of a bit and reversing a bit (from 1 to 0 or from 0 to 1).



## 8.2 Data Hierarchy (cont.)

- ▶ Programming with data in the low-level form of bits is cumbersome.
- ▶ It's preferable to program with data in forms such as **decimal digits** (0–9), **letters** (A–Z and a–z) and **special symbols** (e.g., \$, @, %, &, \* and many others).
- ▶ Digits, letters and special symbols are referred to as **characters**.
- ▶ The set of all characters used to write programs and represent data items on a particular computer is called that computer's **character set**.
- ▶ Every character in a computer's character set is represented as a pattern of 1s and 0s.
- ▶ **Bytes** are composed of eight bits.



## 8.2 Data Hierarchy (cont.)

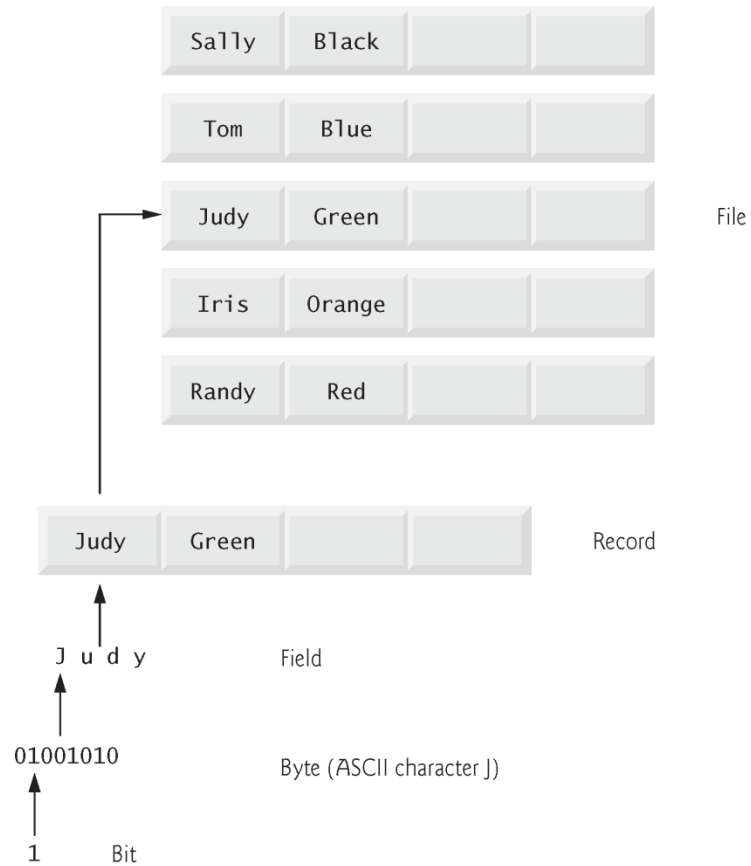
- ▶ You create programs and data items with characters; computers manipulate and process these characters as patterns of bits.
- ▶ Each `char` typically occupies one byte.
- ▶ C++ also provides data type `wchar_t`, which can occupy more than one byte
  - to support larger character sets, such as the [Unicode® character set](#); for more information on Unicode®, visit [www.unicode.org](http://www.unicode.org)



## 8.2 Data Hierarchy (cont.)

- ▶ Just as characters are composed of bits, **fields** are composed of characters.
- ▶ A field is a group of characters that conveys some meaning.
  - For example, a field consisting of uppercase and lowercase letters can represent a person's name.
- ▶ Data items processed by computers form a **data hierarchy** (Fig. 8.1), in which data items become larger and more complex in structure as we progress from bits, to characters, to fields and to larger data aggregates.





**Fig. 8.1** | Data hierarchy.



## 8.2 Data Hierarchy (cont.)

- ▶ Typically, a **record** (which can be represented as a **class** in C++) is composed of several fields (called data members in C++).
  - Thus, a record is a group of related fields.
- ▶ A file is a group of related records.
- ▶ To facilitate retrieving specific records from a file, at least one field in each record is chosen as a **record key**.
- ▶ A record key identifies a record as belonging to a particular person or entity and distinguishes that record from all others.



## 8.2 Data Hierarchy (cont.)

- ▶ There are many ways of organizing records in a file.
- ▶ A common type of organization is called a **sequential file**, in which records typically are stored in order by a record-key field.
- ▶ Most businesses use many different files to store data.
- ▶ A group of related files often are stored in a **database**.
- ▶ A collection of programs designed to create and manage databases is called a **database management system (DBMS)**.



## 8.3 Files and Streams

- ▶ C++ views each file as a sequence of bytes (Fig. 8.2).
- ▶ Each file ends either with an **end-of-file marker** or at a specific byte number recorded in an operating-system-maintained, administrative data structure.
- ▶ When a file is opened, an object is created, and a stream is associated with the object.
- ▶ In Chapter 15, we saw that objects `cin`, `cout`, `cerr` and `clog` are created when `<iostream>` is included.
- ▶ The streams associated with these objects provide communication channels between a program and a particular file or device.



**Fig. 8.2** | C++'s view of a file of  $n$  bytes.



## 8.3 Files and Streams (cont.)

- ▶ To perform file processing in C++, header files `<iostream>` and `<fstream>` must be included.



## 8.4 Creating a Sequential File

- ▶ C++ imposes no structure on a file.
- ▶ Thus, a concept like that of a “record” does not exist in a C++ file.
- ▶ You must structure files to meet the application’s requirements.
- ▶ Figure 8.3 creates a sequential file that might be used in an accounts-receivable system to help manage the money owed by a company’s credit clients.
- ▶ For each client, the program obtains the client’s account number, name and balance (i.e., the amount the client owes the company for goods and services received in the past).
- ▶ The data obtained for each client constitutes a record for that client.
- ▶ The account number serves as the record key.
- ▶ This program assumes the user enters the records in account number order.
  - In a comprehensive accounts receivable system, a sorting capability would be provided to eliminate this restriction.



```
1 // Fig. 8.3: Fig08_03.cpp
2 // Create a sequential file.
3 #include <iostream>
4 #include <string>
5 #include <fstream> // file stream
6 #include <cstdlib>
7 using namespace std;
8
9 int main()
10 {
11     // ofstream constructor opens file
12     ofstream outClientFile( "clients.txt", ios::out );
13
14     // exit program if unable to create file
15     if ( !outClientFile ) // overloaded ! operator
16     {
17         cerr << "File could not be opened" << endl;
18         exit( 1 );
19     } // end if
20
21     cout << "Enter the account, name, and balance." << endl
22         << "Enter end-of-file to end input.\n? ";
23 }
```

**Fig. 8.3** | Creating a sequential file. (Part I of 2.)





```
24  int account; // customer's account number
25  string name; //customer's name
26  double balance; // amount of money customer owes company
27
28  // read account, name and balance from cin, then place in file
29  while ( cin >> account >> name >> balance )
30  {
31      outClientFile << account << ' ' << name << ' ' << balance << endl;
32      cout << "? ";
33  } // end while
34  } // end main
```

```
Enter the account, name, and balance.
Enter end-of-file to end input.
? 100 Jones 24.98
? 200 Doe 345.67
? 300 White 0.00
? 400 Stone -42.16
? 500 Rich 224.62
? ^Z
```

**Fig. 8.3** | Creating a sequential file. (Part 2 of 2.)



## 8.4 Creating a Sequential File (cont.)

- ▶ In Fig. 8.3, the file is to be opened for output, so an `ofstream` object is created.
- ▶ Two arguments are passed to the object's constructor—the `filename` and the `file-open mode` (line 12).
- ▶ For an `ofstream` object, the file-open mode can be either `ios::out` to output data to a file or `ios::app` to append data to the end of a file (without modifying any data already in the file).
- ▶ Existing files opened with mode `ios::out` are `truncated`—all data in the file is discarded.
- ▶ If the specified file does not yet exist, then the `ofstream` object creates the file, using that filename.
- ▶ The `ofstream` constructor opens the file—this establishes a “line of communication” with the file.
- ▶ By default, `ofstream` objects are opened for output, so the open mode is not required in the constructor call.
- ▶ Figure 8.4 lists the file-open modes.



## Common Programming Error 8.1

*Use caution when opening an existing file for output (`ios::out`), especially when you want to preserve the file's contents, which will be discarded without warning.*



Mode	Description
<code>ios::app</code>	Append all output to the end of the file.
<code>ios::ate</code>	Open a file for output and move to the end of the file (normally used to append data to a file). Data can be written anywhere in the file.
<code>ios::in</code>	Open a file for input.
<code>ios::out</code>	Open a file for output.
<code>ios::trunc</code>	Discard the file's contents (this also is the default action for <code>ios::out</code> ).
<code>ios::binary</code>	Open a file for binary (i.e., nontext) input or output.

**Fig. 8.4** | File open modes.



## 8.4 Creating a Sequential File (cont.)

- ▶ An `ofstream` object can be created without opening a specific file—a file can be attached to the object later.
- ▶ For example, the statement
  - `ofstream outClientFile;`
- ▶ creates an `ofstream` object named `outClientFile`.
- ▶ The `ofstream` member function `open` opens a file and attaches it to an existing `ofstream` object as follows:
  - `outClientFile.open("clients.txt", ios::out);`



## **Common Programming Error 8.2**

*Not opening a file before attempting to reference it in a program will result in an error.*



## 8.4 Creating a Sequential File (cont.)

- ▶ After creating an `ofstream` object and attempting to open it, the program tests whether the open operation was successful.
- ▶ The condition in the `if` statement in lines 15–19 returns true if the `open` operation failed.
- ▶ Some possible errors are
  - attempting to open a nonexistent file for reading,
  - attempting to open a file for reading or writing without permission, and
  - opening a file for writing when no disk space is available.



## 8.4 Creating a Sequential File (cont.)

- ▶ Function `exit` terminates a program.
  - The argument to `exit` is returned to the environment from which the program was invoked.
  - Argument `0` indicates that the program terminated normally; any other value indicates that the program terminated due to an error.
  - The calling environment (most likely the operating system) uses the value returned by `exit` to respond appropriately to the error.





## 8.4 Creating a Sequential File (cont.)

- ▶ The `while` statement of lines 29–33 inputs each set of data from the keyboard.
- ▶ The user enters the end-of-file key combination to inform the program to process no additional data—this sets the “end-of-file indicator” in the `cin` object.
- ▶ When the end-of-file indicator is set, the `while` condition becomes false terminating the `while` statement.
- ▶ Figure 8.5 lists the keyboard combinations for entering end-of-file for various computer systems.
- ▶ Later in the chapter, we’ll use the `eof` member function to test for end-of-file in an input file.



Computer system	Keyboard combination
UNIX/Linux/Mac OS X	<Ctrl-d> (on a line by itself)
Microsoft Windows	<Ctrl-z> (sometimes followed by pressing <i>Enter</i> )
VAX (VMS)	<Ctrl-z>

**Fig. 8.5** | End-of-file key combinations for various popular computer systems.



## 8.4 Creating a Sequential File (cont.)

- ▶ Line 31 writes a set of data to the file `clients.txt`, using the stream insertion operator `<<` and the `outClientFile` object associated with the file at the beginning of the program.
- ▶ The data may be retrieved by a program designed to read the file (see Section 8.5).
- ▶ The file created in Fig. 8.3 is simply a text file, so it can be viewed by any text editor.



## 8.4 Creating a Sequential File (cont.)

- ▶ Once the user enters the end-of-file indicator, `main` terminates.
- ▶ This implicitly invokes `outClientFile`'s destructor, which closes the `clients.txt` file.
- ▶ You also can close the `ofstream` object explicitly, using member function `close` in the statement



### Good Programming Practice 8.1

*Open a file for input only (using `ios::in`) if the file's contents should not be modified. This prevents unintentional modification of the file's contents and is an example of the principle of least privilege.*



## 8.5 Reading Data from a Sequential File

- ▶ Files store data so it may be retrieved for processing when needed.
- ▶ In this section, we discuss how to read data sequentially from a file.
- ▶ Figure 8.6 reads records from the `clients.txt` file that we created using the program of Fig. 8.3 and displays the contents of these records.

## 8.5 Reading Data from a Sequential File (cont.)



- ▶ Creating an `ifstream` object opens a file for input.
- ▶ The `ifstream` constructor can receive the filename and the file open mode as arguments.
- ▶ Line 15 creates an `ifstream` object called `inClientFile` and associates it with the `clients.txt` file.
- ▶ The arguments in parentheses are passed to the `ifstream` constructor function, which opens the file and establishes a “line of communication” with the file.



```
1  // Fig. 8.6: Fig08_06.cpp
2  // Reading and printing a sequential file.
3  #include <iostream>
4  #include <fstream> // file stream
5  #include <iomanip>
6  #include <string>
7  #include <cstdlib>
8  using namespace std;
9
10 void outputLine( int, const string, double ); // prototype
11
12 int main()
13 {
14     // ifstream constructor opens the file
15     ifstream inClientFile( "clients.txt", ios::in );
16
17     // exit program if ifstream could not open file
18     if ( !inClientFile )
19     {
20         cerr << "File could not be opened" << endl;
21         exit( 1 );
22     } // end if
23
```

**Fig. 8.6** | Reading and printing a sequential file. (Part I of 3.)





```
24     int account; // customer's account number
25     string name; // customer's name
26     double balance; //amount of money customer owes company
27
28     cout << left << setw( 10 ) << "Account" << setw( 13 )
29         << "Name" << "Balance" << endl << fixed << showpoint;
30
31     // display each record in file
32     while ( inClientFile >> account >> name >> balance )
33         outputLine( account, name, balance );
34 } // end main
35
36 // display single record from file
37 void outputLine( int account, const string name, double balance )
38 {
39     cout << left << setw( 10 ) << account << setw( 13 ) << name
40         << setw( 7 ) << setprecision( 2 ) << right << balance << endl;
41 } // end function outputLine
```

**Fig. 8.6** | Reading and printing a sequential file. (Part 2 of 3.)



Account	Name	Balance
100	Jones	24.98
200	Doe	345.67
300	White	0.00
400	Stone	-42.16
500	Rich	224.62

**Fig. 8.6** | Reading and printing a sequential file. (Part 3 of 3.)

## 8.5 Reading Data from a Sequential File (cont.)



- ▶ Objects of class `ifstream` are opened for input by default, so to open `clients.txt` for input we could have used the statement
  - `ifstream inClientFile( "clients.txt" );`
- ▶ Just as with an `ofstream` object, an `ifstream` object can be created without opening a specific file, because a file can be attached to it later.
- ▶ Each time line 32 executes, it reads another record from the file into the variables `account`, `name` and `balance`.
- ▶ When the end of file has been reached, the `while` condition returns `false`), terminating the `while` statement and the program; this causes the `ifstream` destructor function to run, closing the file.

## 8.5 Reading Data from a Sequential File (cont.)



- ▶ To retrieve data sequentially from a file, programs normally start reading from the beginning of the file and read all the data consecutively until the desired data is found.
- ▶ It might be necessary to process the file sequentially several times (from the beginning of the file) during program execution.
- ▶ Both `istream` and `ostream` provide member functions for repositioning the **file-position pointer** (the byte number of the next byte in the file to be read or written).
  - `seekg` (“seek get”) for `istream`
  - `seekp` (“seek put”) for `ostream`

## 8.5 Reading Data from a Sequential File (cont.)



- ▶ Each `istream` object has a “get pointer,” which indicates the byte number in the file from which the next input is to occur, and each `ostream` object has a “put pointer,” which indicates the byte number in the file at which the next output should be placed.
- ▶ The statement
  - `inClientFile.seekg( 0 );`repositions the file-position pointer to the beginning of the file (location 0) attached to `inClientFile`.
- ▶ The argument to `seekg` normally is a `long` integer.

## 8.5 Reading Data from a Sequential File (cont.)



- ▶ A second argument can be specified to indicate the **seek direction**, which can be
  - `ios::beg` (the default) for positioning relative to the beginning of a stream,
  - `ios::cur` for positioning relative to the current position in a stream or
  - `ios::end` for positioning relative to the end of a stream
- ▶ The file-position pointer is an integer value that specifies the location in the file as a number of bytes from the file's starting location (this is also referred to as the **offset** from the beginning of the file).

## 8.5 Reading Data from a Sequential File (cont.)



- ▶ Some examples of positioning the “get” file-position pointer are
  - `// position to the nth byte of fileObject (assumes ios::beg)`  
`fileObject.seekg( n );`
  - `// position n bytes forward in fileObject`  
`fileObject.seekg( n, ios::cur );`
  - `// position n bytes back from end of fileObject`  
`fileObject.seekg( n, ios::end );`
  - `// position at end of fileObject`  
`fileObject.seekg( 0, ios::end );`
- ▶ The same operations can be performed using `ostream` member function `seekp`.

## 8.5 Reading Data from a Sequential File (cont.)



- ▶ Member functions `tellg` and `tellp` are provided to return the current locations of the “get” and “put” pointers, respectively.
- ▶ Figure 8.7 enables a credit manager to display the account information for those customers with
  - zero balances (i.e., customers who do not owe the company any money),
  - credit (negative) balances (i.e., customers to whom the company owes money), and
  - debit (positive) balances (i.e., customers who owe the company money for goods and services received in the past)





```
1 // Fig. 8.7: Fig08_08.cpp
2 // Credit inquiry program.
3 #include <iostream>
4 #include <fstream>
5 #include <iomanip>
6 #include <string>
7 #include <cstdlib>
8 using namespace std;
9
10 enum RequestType { ZERO_BALANCE = 1, CREDIT_BALANCE, DEBIT_BALANCE, END };
11 int getRequest();
12 bool shouldDisplay( int, double );
13 void outputLine( int, const string, double );
14
15 int main()
16 {
17     // ifstream constructor opens the file
18     ifstream inClientFile( "clients.txt", ios::in );
19
20     // exit program if ifstream could not open file
21     if ( !inClientFile )
22     {
23         cerr << "File could not be opened" << endl;
24         exit( 1 );
25     } // end if
```

**Fig. 8.7** | Credit inquiry program. (Part I of 7.)



```
26
27  int request; // request type: zero, credit or debit balance
28  int account; // customer's account number
29  string name; // customer's name
30  double balance; // amount of money customer owes company
31
32  // get user's request (e.g., zero, credit or debit balance)
33  request = getRequest();
34
35  // process user's request
36  while ( request != END )
37  {
38      switch ( request )
39      {
40          case ZERO_BALANCE:
41              cout << "\nAccounts with zero balances:\n";
42              break;
43          case CREDIT_BALANCE:
44              cout << "\nAccounts with credit balances:\n";
45              break;
46          case DEBIT_BALANCE:
47              cout << "\nAccounts with debit balances:\n";
48              break;
49      } // end switch
```

**Fig. 8.7** | Credit inquiry program. (Part 2 of 7.)



```
50
51 // read account, name and balance from file
52 inClientFile >> account >> name >> balance;
53
54 // display file contents (until eof)
55 while ( !inClientFile.eof() )
56 {
57     // display record
58     if ( shouldDisplay( request, balance ) )
59         outputLine( account, name, balance );
60
61     // read account, name and balance from file
62     inClientFile >> account >> name >> balance;
63 } // end inner while
64
65 inClientFile.clear(); // reset eof for next input
66 inClientFile.seekg( 0 ); // reposition to beginning of file
67 request = getRequest(); // get additional request from user
68 } // end outer while
69
70 cout << "End of run." << endl;
71 } // end main
72
```

**Fig. 8.7** | Credit inquiry program. (Part 3 of 7.)



---

```
73 // obtain request from user
74 int getRequest()
75 {
76     int request; // request from user
77
78     // display request options
79     cout << "\nEnter request" << endl
80         << " 1 - List accounts with zero balances" << endl
81         << " 2 - List accounts with credit balances" << endl
82         << " 3 - List accounts with debit balances" << endl
83         << " 4 - End of run" << fixed << showpoint;
84
85     do // input user request
86     {
87         cout << "\n? ";
88         cin >> request;
89     } while ( request < ZERO_BALANCE && request > END );
90
91     return request;
92 } // end function getRequest
93
```

---

**Fig. 8.7** | Credit inquiry program. (Part 4 of 7.)



---

```
94 // determine whether to display given record
95 bool shouldDisplay( int type, double balance )
96 {
97     // determine whether to display zero balances
98     if ( type == ZERO_BALANCE && balance == 0 )
99         return true;
100
101     // determine whether to display credit balances
102     if ( type == CREDIT_BALANCE && balance < 0 )
103         return true;
104
105     // determine whether to display debit balances
106     if ( type == DEBIT_BALANCE && balance > 0 )
107         return true;
108
109     return false;
110 } // end function shouldDisplay
111
```

---

**Fig. 8.7** | Credit inquiry program. (Part 5 of 7.)



```
112 // display single record from file
113 void outputLine( int account, const string name, double balance )
114 {
115     cout << left << setw( 10 ) << account << setw( 13 ) << name
116         << setw( 7 ) << setprecision( 2 ) << right << balance << endl;
117 } // end function outputLine
```

```
Enter request
1 - List accounts with zero balances
2 - List accounts with credit balances
3 - List accounts with debit balances
4 - End of run
? 1
```

```
Accounts with zero balances:
300      White      0.00
```

```
Enter request
1 - List accounts with zero balances
2 - List accounts with credit balances
3 - List accounts with debit balances
4 - End of run
? 2
```

**Fig. 8.7** | Credit inquiry program. (Part 6 of 7.)



Accounts with credit balances:

400	Stone	-42.16
-----	-------	--------

Enter request

- 1 - List accounts with zero balances
- 2 - List accounts with credit balances
- 3 - List accounts with debit balances
- 4 - End of run

? 3

Accounts with debit balances:

100	Jones	24.98
200	Doe	345.67
500	Rich	224.62

Enter request

- 1 - List accounts with zero balances
- 2 - List accounts with credit balances
- 3 - List accounts with debit balances
- 4 - End of run

? 4

End of run.

**Fig. 8.7** | Credit inquiry program. (Part 7 of 7.)



## 8.6 Updating Sequential Files

- ▶ Data that is formatted and written to a sequential file as shown in Section 8.4 cannot be modified without the risk of destroying other data in the file.
- ▶ For example, if the name “white” needs to be changed to “worthington,” the old name cannot be overwritten without corrupting the file.
- ▶ The record for white was written to the file as
  - 300 white 0.00
- ▶ If this record were rewritten beginning at the same location in the file using the longer name, the record would be
  - 300 worthington 0.00
- ▶ The new record contains six more characters than the original record.
- ▶ Therefore, the characters beyond the second “o” in “worthington” would overwrite the beginning of the next sequential record in the file.





## 8.6 Updating Sequential Files (cont.)

- ▶ The problem is that, in the formatted input/output model using the stream insertion operator `<<` and the stream extraction operator `>>`, fields—and hence records—can vary in size.
  - For example, values 7, 14, -117, 2074, and 27383 are all `ints`, which store the same number of “raw data” bytes internally (typically four bytes on today’s popular 32-bit machines).
  - However, these integers become different-sized fields when output as formatted text (character sequences).
  - Therefore, the formatted input/output model usually is not used to update records in place.



## 8.6 Updating Sequential Files (cont.)

- ▶ Such updating can be done, but a bit awkwardly.
- ▶ For example, to make the preceding name change
  - the records before 300 White 0.00 in a sequential file could be copied to a new file
  - the updated record then written to the new file
  - and the records after 300 White 0.00 copied to the new file.
- ▶ This requires processing *every* record in the file to update *only* one record.
- ▶ If many records are being updated in one pass of the file, though, this technique can be acceptable.